

# Building resilient applications on Azure

27. April 2022



# Presenting today



**Ricardo Niepel**

Cloud Solution Architect – Azure App Dev

[ricardo.niepel@microsoft.com](mailto:ricardo.niepel@microsoft.com)



RicardoNiepel



@RicardoNiepel



**Dennis Zielke**

Technical Lead - Cloud Native Apps

[dennis.zielke@microsoft.com](mailto:dennis.zielke@microsoft.com)




denniszielke



@denzielke

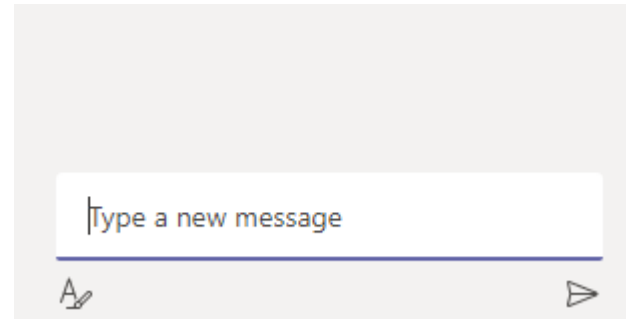
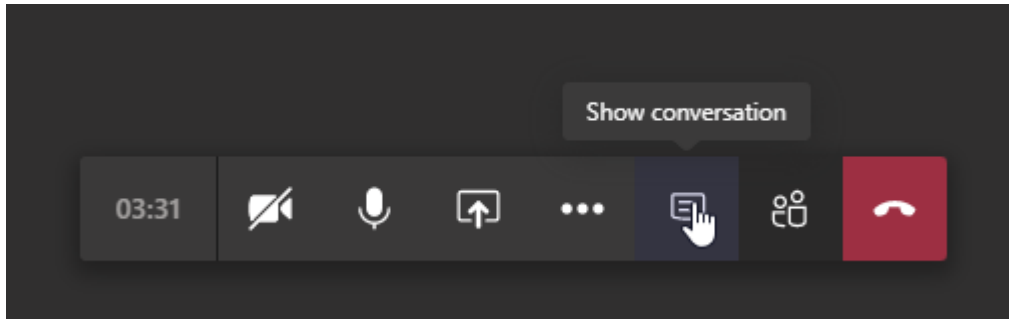
# Organizational Upfront

- Schedule: Monthly, 2nd or 3rd Friday
- If your colleagues have trouble joining this Teams call:  
Use the Web Client (no plug-in) - Google Chrome version >72 or Microsoft Edge (Chromium-based)
- No recordings / if you miss it pls join next time
- Please make sure to mute yourself, too many people



# Questions

If you have any Cloud Native relevant question, please use the Teams chat

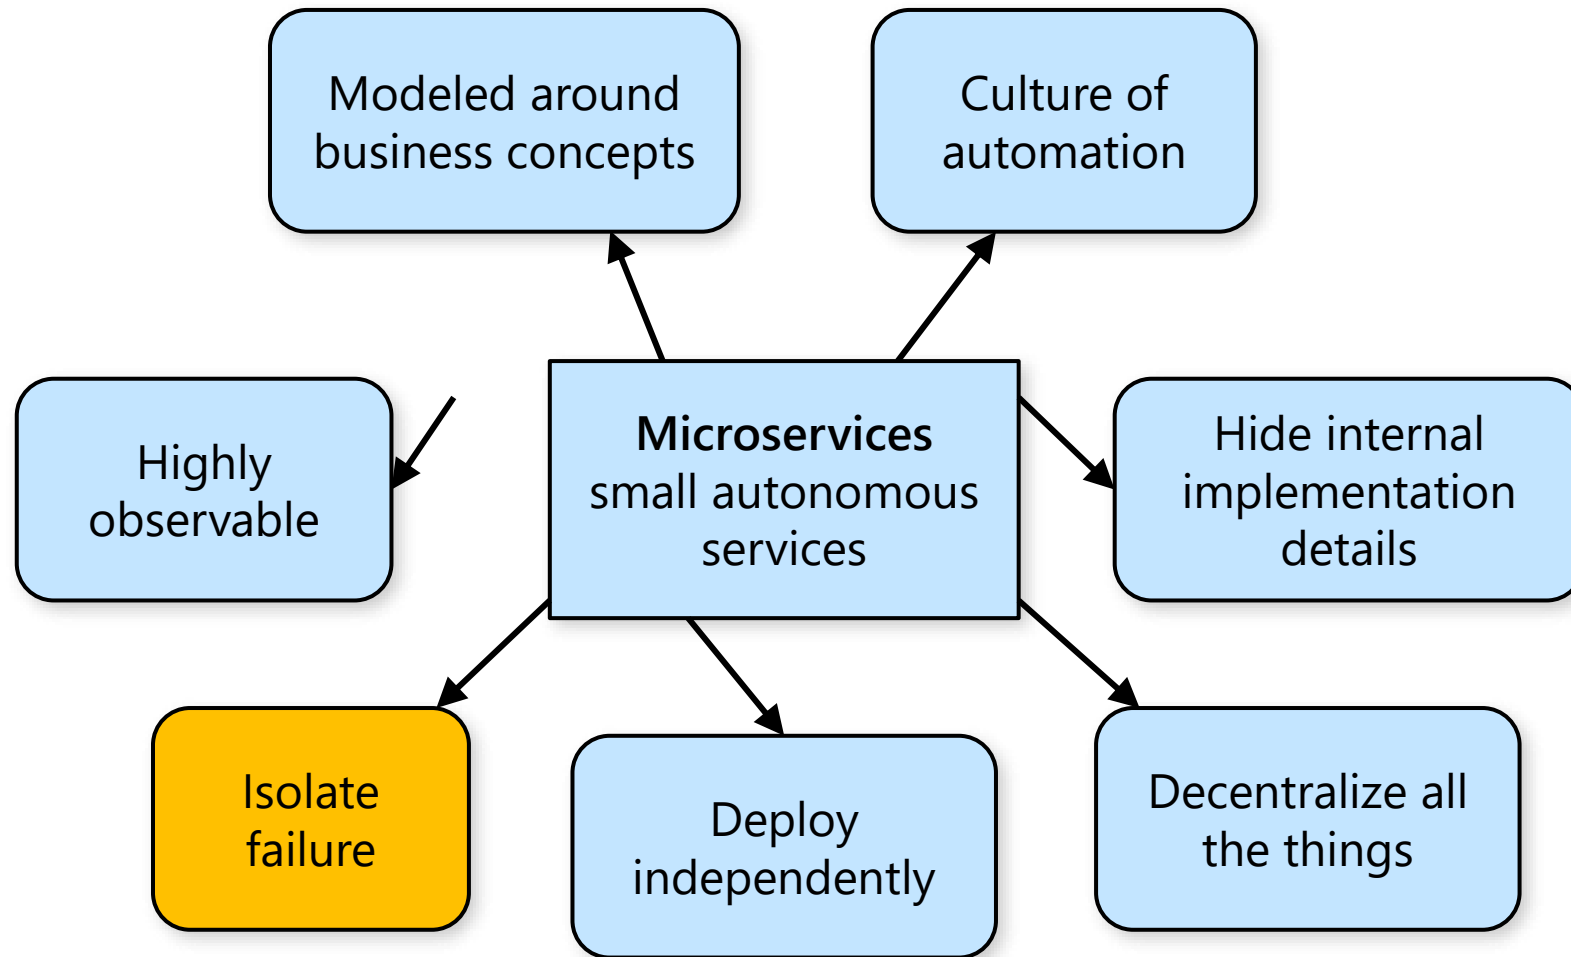


If we have time left at the end, we will pick some questions and answering them live

# Agenda for today

- How can cloud native applications fail?
- How can you build your application resilient to failure?
- How can I implement a resilient application architecture?
- How can you validate your assumptions?

# Principles of microservices



# Reliability and Resilience


**Reliable** workloads are:

**Highly available (HA)** – amount of uptime a system is ready to perform core or essential functions

**Resilient** - able to recover gracefully from failures and continue to function with minimal downtime and data loss

# Views on Failure

## Failure of a component




Focus on: Minimizing risk of failure

Questions to ask:

- How to implement error handling?
- What service SKU to select?
- What redundancy level to choose?

## Failure of a dependency



Focus on: Detecting and isolating failures

Questions to ask:

- How to detect failures?
- How to protect myself?
- How to support recovery?



# Infrastructure failures are rare but can happen

## Zonal failure

Some compute will be completely unavailable

## Load Balancer failure

Load balancing from the outside to some nodes will fail

## Node failure

A compute node will fail completely

## Disk failure

What is the max duration of data loss that is acceptable?

**-> Can be solved by having sufficient available infrastructure resources**

# Partial infrastructure failures are more common

## Health probe unreliability

Service works but health probe do not reflect true state

## Node overload

Service runs but cannot respond in a timely manner due to CPU/ memory/ disk

## Port exhaustion

Service works but cannot connect to its dependencies

## DNS resolution failures

Service works but cannot resolve its dependencies

**-> Needs to be solved by application design to detect and handle accordingly**

# Traits of Resilient Systems

## Handle transient failures

When failure is encountered, retry, queue and orchestrate transactions/operations

## Continue operations during failure

Some part of the system will always be in a failed state.

## Maintain consistent state

Apply compensating logic across multiple data stores or attempt resubmission until operation succeeds.

# Circuit Breaker - *Stop hitting it if it hurts.*

Dependency  
failure

## Solution

Prevent an application from repeatedly trying to execute an operation that is likely to fail and enables an application to detect whether the fault has been resolved.


If the problem appears to have been fixed, the application can try to invoke the operation.

## Considerations

Circuit Breaker isn't just a more elaborate Retry—it's based on the premises that subsequent calls will likely *not* succeed, whereas Retry assume they will likely succeed.

## Beware

Circuit Breaker are inherently stateful.



# Timeout - *Don't wait forever.*

DNS failure

## Solution

Only allow an operation a limited amount of time to complete.

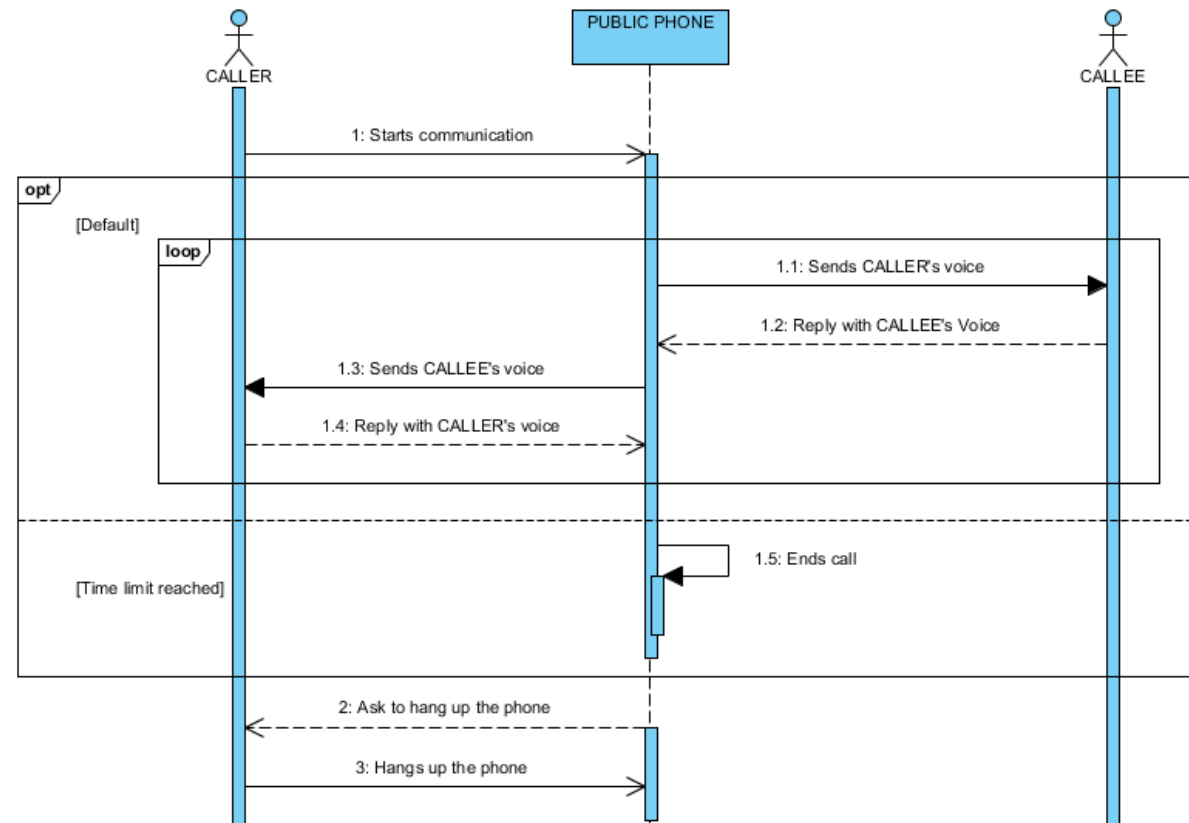
Waiting forever (having no timeout) is a bad design strategy. It leads to the blocking up of threads or connections (itself often a cause of further failure), during a faulting scenario.

## Considerations

May require operation to run on an additional thread, if it's not cancellable, doubling the number of threads you actually require.

## Beware

Easy with cancellable operations. If this isn't possible, Timeouts can become quite tricky.



# Retry - *Maybe it's just a blip!*

Health Probe  
failure

## Solution

Transparently handle failures as if the failure condition will resolve itself.

## Considerations


Only works if you can identify transient errors.

Use exponential back-off to allow for retries to be made initially quickly, but then at progressively longer intervals, to avoid hitting a subsystem with repeated frequent calls if the subsystem may be struggling.

Apply jitter to prevent retries bunching into further spikes of load.

## Beware

Only retry idempotent operations or provide compensating logic.



- 1: Application invokes operation on hosted service. The request fails, and the service host responds with HTTP response code 500 (internal server error).
- 2: Application waits for a short interval and tries again. The request still fails with HTTP response code 500.
- 3: Application waits for a longer interval and tries again. The request succeeds with HTTP response code 200 (OK).

# Throttling - *Enough is enough.*

## Solution

Allow applications to use resources only up to a limit, and then throttle them when this limit is reached.

This can allow the system to continue to function and meet service level agreements, even when an increase in demand places an extreme load on resources.


## Considerations

Often used as a building block for a higher-level larger pattern (e.g. bulkhead).

Rate limiting (i.e. calls/sec) or leasing (i.e. a pool of  $n$  resources) are typical stand-alone versions of Throttling.

## Considerations

Familiarize yourself with the Throttling behaviors of Azure services and APIs!



# Queue Load Leveling - *Keep calm and dispatch a message.*

## Solution

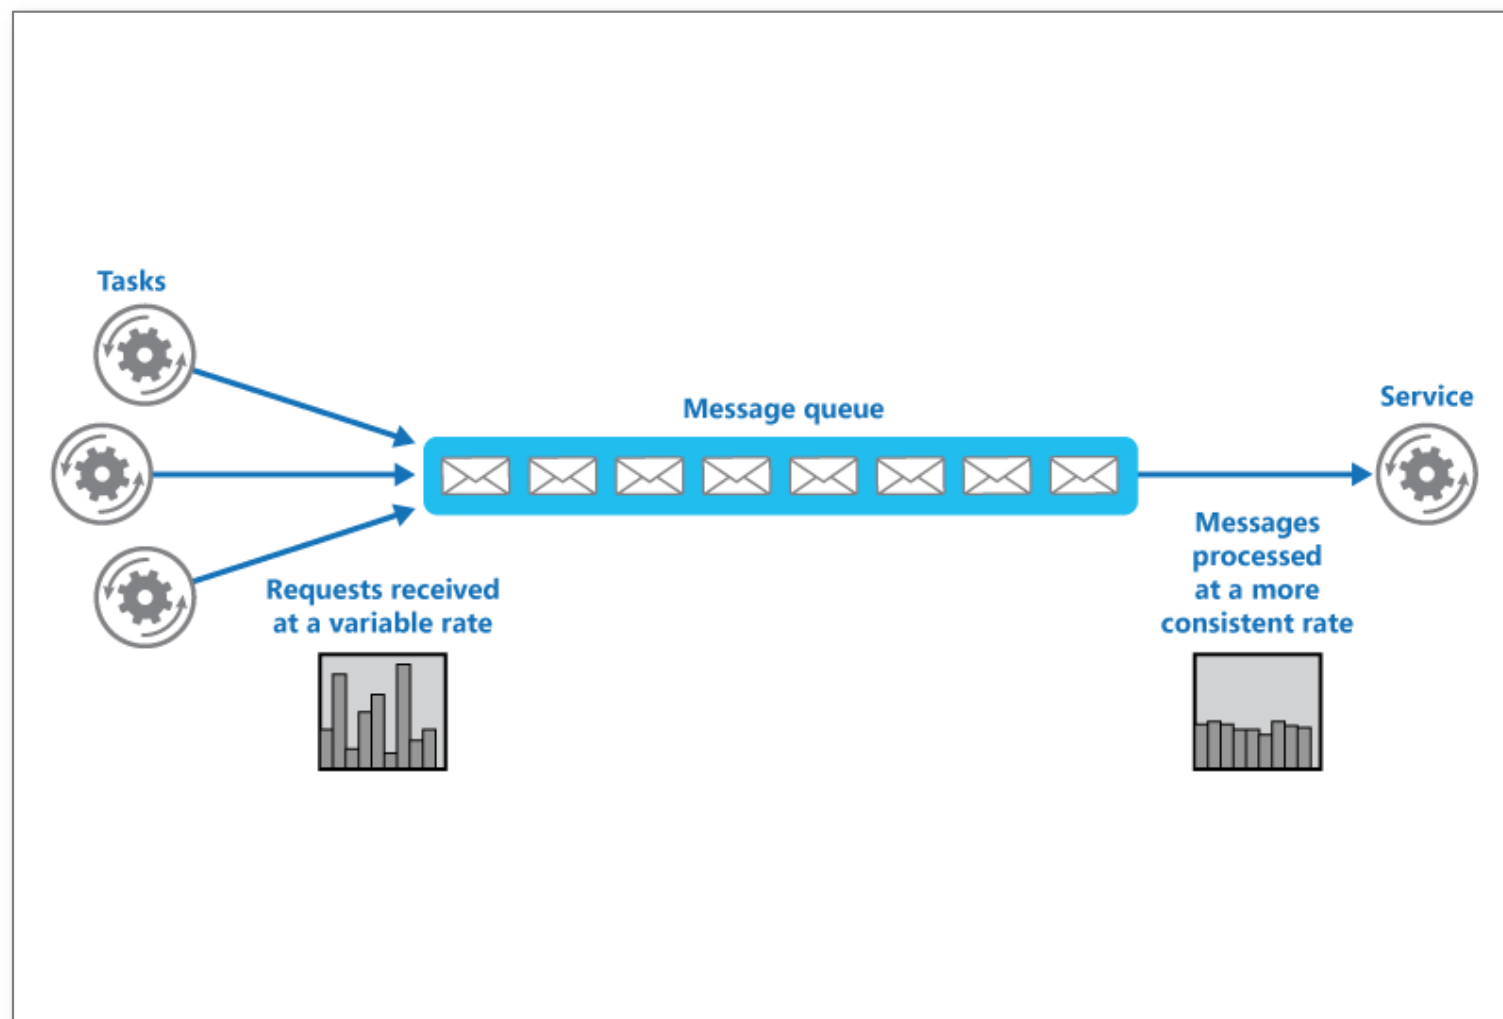
Introduce a queue between the task and the service. The task and the service run asynchronously at their own processing speed.

## Considerations

Requires a separate channel for clients to receive a task's result (i.e. pub/sub, web sockets, etc).

## Beware

Introduces a new dependency (i.e. the queue) that may fail as well.





Thank you.

Slides & FAQ: